

EU AI ACT

HIDDEN CLIENT SCENARIO BRIEF

Consulting Exercise — Elsa's Copy

Case 1

Client need:

A digital investment platform wants to increase user engagement and trading activity by guiding users toward more optimal investment decisions and reducing impulsive behavior during market volatility.

Proposed AI behaviour:

The system will analyze user activity such as trading history, browsing patterns, reaction times, and interaction signals to infer behavioral states like hesitation, fear, or overconfidence. Based on these signals, it will dynamically adjust the interface, including highlighting specific assets, delaying or accelerating certain actions, and nudging users toward or away from trades. The system operates in real time without human review of individual interventions.

People affected:

Retail investors using the platform, whose investment decisions may be influenced by system-driven nudges based on inferred behavioral traits, without visibility into how or why these interventions occur.

Case 2

Client need:

A company responsible for maintaining waterway infrastructure wants to improve the efficiency and consistency of field inspections while reducing manual reporting effort for engineers.

Proposed AI behaviour:

The system will ingest voice recordings, images, and sensor data collected on-site. It will validate pre-identified damage suspicions, classify damage types, estimate severity, and suggest repair measures. It will generate structured inspection reports automatically, which engineers can review, edit, and approve before submitting them into the central system used for maintenance planning and budgeting.

People affected:

Field inspectors and office engineers who rely on the system's outputs to document and assess infrastructure conditions, as well as the broader public who depend on the safety and reliability of the infrastructure being inspected.

Case 3

Client need:

A booking platform for beauty and wellness services wants to improve customer experience and reduce operational workload by automating customer interactions and appointment management.

Proposed AI behaviour:

The system will act as a conversational assistant that helps users find providers, book appointments, reschedule services, and receive recommendations based on preferences and past behavior. It will handle interactions through chat and voice interfaces and complete most actions end-to-end. The system presents itself as a digital assistant but does not explicitly clarify whether responses are generated by AI.

People affected:

Platform users seeking services, who rely on the assistant for recommendations and booking decisions without always knowing whether they are interacting with an automated system or a human.

Case 4

Client need:

A startup developing a productivity tool wants to help teams reduce time spent on meeting follow-ups and improve alignment by automatically summarizing discussions and identifying next steps.

Proposed AI behaviour:

The system will process uploaded meeting recordings and transcripts to generate concise summaries and suggested action items. These outputs will be presented in a structured format and can be edited, approved, or discarded by users before being shared with others. The system functions as a drafting assistant and does not make or enforce decisions.

People affected:

Employees and team members using the tool to document meetings and coordinate tasks, whose work outputs may be influenced by AI-generated summaries that are subject to human review before use.

2. Private Answer Key

Case	Classification	Why?
Case 1 Digital Investment	PROHIBITED — Art. 5(1)(f)	The system infers emotional/behavioral states and uses them in real time to steer user decisions without awareness. This introduces covert influence over financial decisions, limiting user autonomy and making outcomes non-transparent and non-contestable.
Case 2 Water Infrastructure	HIGH-RISK — Annex III, Area 5	The system supports decisions related to infrastructure condition and repair, which directly impacts public safety and asset integrity. Although human review exists, the AI materially influences high-stakes technical assessments, triggering high-risk classification.
Case 3 Booking Platform	LIMITED RISK — Art. 50 transparency	The system interacts directly with users in a human-like manner without clearly disclosing that it is AI. This creates a transparency gap, as users may rely on recommendations without understanding they are generated by an automated system.
Case 4 Meeting Summarisation tool	MINIMAL RISK — no Annex III trigger	The system generates summaries and suggestions but does not make decisions or impact rights. Outputs are reviewed and controlled by users, and the tool operates as a productivity aid with limited external or individual impact.

Case A

InvestmentPlatform

PROHIBITED (likely)

Article 5(1)(a) + Article 5(1)(b) EU AI Act

Likely category

PROHIBITED — Art. 5(1)(a) and Art. 5(1)(b)

Why this is your first-pass call

Two prohibited elements are likely triggered:

(1) Art. 5(1)(a)/(b) — the system infers emotional and behavioral states (fear, hesitation, overconfidence) and uses them in real time to steer user decisions through interface manipulation (nudging, delaying actions, highlighting assets). This materially distorts user behavior without their awareness.

(2) The influence occurs in a **financial decision-making context**, increasing the risk of harm and reducing user autonomy.

The key issue is not personalization itself, but **covert behavioral manipulation at decision time**, with no transparency, no user control, and no human oversight. The system is designed to influence outcomes in a way users cannot detect or contest.

Proposed AI architecture

User interacts with trading platform → system collects behavioral signals (trading history, browsing patterns, reaction times) → AI model infers emotional/behavioral state → interface dynamically adapts (highlight assets, delay/accelerate actions, nudge decisions) → user decision influenced in real time without disclosure → no human review or override.

Provider / deployer / vendor

Provider: Internal ML system or behavioral analytics vendor

Deployer: Investment platform

Vendor: Potential third-party analytics or data enrichment providers

Required obligations or controls

No controls can make this lawful as designed. Redesign required:

- Remove real-time behavioral manipulation layer
- Avoid inference of emotional states for decision steering
- Ensure transparency of recommendation logic
- Shift to **user-controlled, explainable recommendation systems**
- Introduce optional, non-intrusive guidance rather than covert nudging

Decision

✗ DENY AND REDESIGN

Case B

InfraCloud Inspection Assistant

HIGH-RISK

Annex III — Critical infrastructure / safety-related systems

Likely category

HIGH-RISK — Annex III (infrastructure safety and maintenance)

Why this is your first-pass call

The system supports inspection and maintenance decisions for infrastructure assets. These decisions directly impact **public safety, asset integrity, and financial planning**.

Even though the AI does not make final decisions autonomously, it:

- classifies damage
- estimates severity
- suggests repair measures

This constitutes a **material influence on high-stakes technical decisions**, placing it within Annex III scope.

The presence of human review is positive but does not remove the classification — it instead defines the required controls.

Proposed AI architecture

Inspector collects field data (voice, images, sensor input) → AI processes inputs → validates damage suspicions + classifies damage + estimates severity → generates structured report + repair suggestions → inspector reviews, edits, and confirms → data submitted to central system → used for planning and budgeting decisions.

Provider / deployer / vendor

Provider: AI components (STT, CV, ML models — internal or third-party)

Deployer: Infrastructure operator / engineering company

Vendor: Possible AI service providers (vision, speech, LLM APIs)

Required obligations or controls

System can be compliant with strong controls:

- Mandatory human validation before submission
- Confidence scores and explainability of outputs
- Full audit trail of AI suggestions and human edits
- Monitoring of model accuracy and failure cases
- Clear fallback to manual workflows
- Training for inspectors on system limitations

Decision

 **APPROVE WITH CONTROLS**

Case C

Rumio AI Assistant

LIMITED RISK

Article 50 EU AI Act

Likely category

LIMITED RISK — Art. 50 transparency obligation

Why this is your first-pass call

The system interacts directly with users and influences their choices (booking, provider selection), but:

- does not determine access to essential services
- does not operate in a high-risk domain

The main issue is **lack of transparency**:

- users are not clearly informed they are interacting with AI
- the assistant presents itself in a human-like way

This creates a risk of **misleading interaction**, which falls under Article 50.

Proposed AI architecture

User interacts via chat/voice → AI assistant processes request → recommends providers, books appointments, handles rescheduling → executes actions end-to-end → no explicit disclosure of AI → optional or missing human escalation.

Provider / deployer / vendor

Provider: LLM / AI assistant vendor

Deployer: Rumio platform

Vendor: Third-party AI APIs (LLM, voice, etc.)

Required obligations or controls

System is compliant with minimal adjustments:

- Clear disclosure at start of interaction (“You are chatting with an AI assistant”)
- Accessible human escalation option
- Logging of interactions for traceability
- Guardrails for recommendations (avoid misleading or biased suggestions)

Decision

 **APPROVE WITH CONTROLS**

Case D

Meeting Summarizer

MINIMAL RISK

No Annex III trigger

Likely category

MINIMAL RISK — no Annex III trigger

Why this is your first-pass call

The system operates as a **drafting and productivity tool**, not a decision-maker.

- It generates summaries and suggested actions
- Outputs are always reviewed and controlled by users
- It does not affect rights, access to services, or critical decisions

The impact is limited to **internal workflow efficiency**, with full human control retained.

Proposed AI architecture

User uploads meeting recording/transcript → AI generates summary + action items → structured output presented → user edits/approves/discards → final version shared manually.

Provider / deployer / vendor

Provider: LLM or summarization model

Deployer: Startup / internal users

Vendor: AI API providers (if external models used)

Required obligations or controls

No mandatory AI Act obligations. Recommended good practices:

- Human review before sharing outputs
- Version control of generated summaries
- Internal AI inventory documentation
- Data handling safeguards if recordings include personal data

Decision

 APPROVE